

An efficient way of finding all fixed points in shallow (low-rank) piece-wise linear recurrent neural networks

Matthijs Pals, University of Tuebingen

September 21, 2023

1 Introduction

Piece-wise linear recurrent neural networks (PLRNNs) have been shown to be able emulate a wide range of dynamical systems, all while being interpretable - all fixed points can be solved for (given a potentially high computational cost). Recently a 'low-rank' variant of the piece-wise model was proposed (shPLRNN; eq. 12 in [1]):

$$\mathbf{z}_t = \mathbf{A}\mathbf{z}_{t-1} + \mathbf{V}^\top \text{ReLU}(\mathbf{U}\mathbf{z}_{t-1} - \mathbf{h}_\mathbf{x}) - \mathbf{h}_\mathbf{z} \quad (1)$$

with $\mathbf{z}_t \in \mathcal{R}^R$ and $\mathbf{V}, \mathbf{U} \in \mathcal{R}^{N \times R}$. To find all fixed points naively, we need to solve 2^N inverses of $R \times R$ matrices. (eq. 39 in [1])

2 Main result

We can obtain all fixed points of a shPLRNN, by solving $\sum_{r=0}^R \binom{N}{r}$ inverses of $R \times R$ matrices.

3 Argument

3.1 Obtaining fixed points in PLRNNs

First, following [1]. To find all fixed points, start by redefining ReLU by introducing a matrix:

$$\mathbf{D} = \begin{bmatrix} \mathbf{d}_1 & & & \\ & \mathbf{d}_2 & & \\ & & \ddots & \\ & & & \mathbf{d}_N \end{bmatrix} \quad (2)$$

with $\mathbf{d}_i = \begin{cases} 1, & \text{if } \mathbf{x}_i > \mathbf{b}_i \\ 0, & \text{otherwise} \end{cases}$

Then

$$\mathbf{z}_t = \mathbf{A}\mathbf{z}_{t-1} + \mathbf{V}^\top \mathbf{D} \mathbf{U} \mathbf{z}_{t-1} - \mathbf{V}^\top \mathbf{D} \mathbf{h}_\mathbf{x} - \mathbf{h}_\mathbf{z} \quad (3)$$

Next, for all all 2^N configurations $\mathbf{D}^{(j)}$ of \mathbf{D} , solve:

$$\mathbf{z}^* = (\mathbf{A} + \mathbf{V}^\top \mathbf{D}^{(j)} \mathbf{U} - \mathbf{I})^{-1} (\mathbf{V}^\top \mathbf{D}^{(j)} \mathbf{h}_\mathbf{x} + \mathbf{h}_\mathbf{z}) \quad (4)$$

for \mathbf{z}^* , and check whether the obtained \mathbf{z}^* is consistent with the assumed $\mathbf{D}^{(j)}$, if so, you found a fixed point of the shPLRNN.




Figure 1: **a.** We can only reach a subset of the subspaces corresponding to each \mathbf{D} . **b.** Each neuron breaks up the space spanned by \mathbf{u} with a hyperplane (point in the 1D case) **c.** All subspaces of a randomly initialised network with $N=8$, $R=2$.

3.2 Observation 1: a shPLRNN is a low-rank version of a standard PLRNN

Introduce $\mathbf{x}_t \in \mathcal{R}^N$, such that \mathbf{z}_t can be seen as the projection of \mathbf{x}_t on \mathbf{U} : $\mathbf{x}_t = \mathbf{U}\mathbf{z}_t$, $\mathbf{z}_t = (\mathbf{U}^\top \mathbf{U})^{-1} \mathbf{U}^\top \mathbf{x}_t$. Furthermore, introduce $\mathbf{W} = \mathbf{U}\mathbf{V}^\top$ and $\mathbf{A}_\mathbf{x} = \mathbf{U}\mathbf{A}(\mathbf{U}^\top \mathbf{U})^{-1} \mathbf{U}^\top$

Then we can rewrite equation 1 as:

$$\mathbf{x}_t = \mathbf{A}_\mathbf{x} \mathbf{x}_{t-1} + \mathbf{W} \text{ReLU}(\mathbf{x}_{t-1} - \mathbf{h}_\mathbf{x}) - \mathbf{U} \mathbf{h}_\mathbf{z} \quad (5)$$

a 'standard' PLRNN - however \mathbf{W} and $\mathbf{A}_\mathbf{x}$ are low-rank matrices, with span columns of \mathbf{U} . As a consequence \mathbf{x}_t is constrained to an R dimensional linear subspace (see also work by e.g. Ostojic' lab).

3.3 Observation 2: not all configurations of \mathbf{D} can be reached

To get some intuition, consider $R = 1, N = 2$. The phase space of the model can be broken up in 4 subspaces, each corresponding to one configuration of \mathbf{D} (Figure 1a). The dynamics are however confined to the line spanned by the vector \mathbf{u} (with coordinate z), which can reach at most 3 of the subspaces (red line in Figure 1a).

We can solve for the points where \mathbf{u} crosses a border between subspaces. E.g. for border given by $\mathbf{x}_1 = \mathbf{b}_1$, we have $\begin{bmatrix} \mathbf{x}_1 \\ \mathbf{x}_2 \end{bmatrix} = z \begin{bmatrix} \mathbf{u}_1 \\ \mathbf{u}_2 \end{bmatrix} = \begin{bmatrix} \mathbf{b}_1 \\ 0 \end{bmatrix}$, thus this border crosses \mathbf{u} at $z = \frac{\mathbf{b}_1}{\mathbf{u}_1}$ (see Figure 1b)

In general, for networks with $R = 1$, each neuron adds 1 point to the line along which the dynamics flow (at $z = \frac{\mathbf{b}_i}{\mathbf{u}_i}$ for the i 'th neuron), meaning we can reach at most $N + 1$ subspaces / configurations of \mathbf{D} . We only have to solve $N + 1$ 'matrix' inverses to find all inverses!

For $R = 2$ dynamics are confined to a plane, spanned by the two columns of \mathbf{U} , and each neuron adds a line to this plane (Figure 1c contains an example with $N=8$). The amount of configurations of \mathbf{D} we can reach is thus equal to the question, in how many surfaces can we divide a plane with N lines, which is $\frac{1}{2}(N^2 + N) + 1$ — a big improvement over 2^N .

Generalising this to arbitrary R , we have to answer the question, in how many subspaces can we divide an R dimensional space with N hyperplanes, which has the solution $\sum_{r=0}^R \binom{N}{r}$

3.4 Finding all reachable configurations of \mathbf{D}

We so far showed that we only need to consider $\sum_{r=0}^R \binom{N}{r}$ configurations of \mathbf{D} , but how do we find these configurations?. First we can solve for all cross sections of hyperplanes (see previous section for $r=1$). To do this, we have to solve $\binom{N}{r}$ systems of linear equations.

From each cross sections we can reach 2^R possible subregions by moving epsilon away from the cross section. We can find the corresponding \mathbf{D} matrices of each of those subsection as follows. First compute $\mathbf{x} = \mathbf{U}\mathbf{z}$, corresponding to the cross section we are currently considering. Next calculate all the \mathbf{D} s corresponding to the found \mathbf{x} — for elements \mathbf{x}_i of \mathbf{x} which are equal to \mathbf{b}_i , \mathbf{D}_i can be either 1, or 0. Now we just make a list of all possible \mathbf{D} s that we find this way, and finally remove any duplicates.

References

- [1] Florian Hess, Zahra Monfared, Manuel Brenner, and Daniel Durstewitz. Generalized teacher forcing for learning chaotic dynamics, 2023.